Himanshu Singh

+91 9129491699 | dev88himanshu@gmail.com

[LinkedIn](#) ↗ | **[Portfolio](#)** ↗

PROFESSIONAL SUMMARY

Azure Certified Data Engineer with **4+ years of experience** designing and optimizing **data pipelines, lakehouse architectures**, and **data warehouse solutions** across **Azure** and **on-premise Hadoop ecosystems**. Skilled in **ADF, Databricks, Spark, Hive, SQL, and Python**, delivering up to **60% performance improvements** and ensuring **data reliability at scale** for **Fortune 100 Healthcare and Retail (Pharmacy) clients**.

TECHNICAL SKILLS

Cloud & Data Engineering:	ADF, Databricks, Azure Synapse, ADLS/Blob Storage, Data Lakehouse
Big Data Ecosystem:	Spark (PySpark, RDD, Streaming), Hadoop (Hive, Impala, Sqoop, MapReduce, YARN)
Pipelines & Warehousing:	ETL/ELT, Data Modelling, Data Warehousing OLAP/OLTP
Programming:	Python, SQL, Shell Scripting
Orchestration & Workflow:	Tivoli Workload Scheduler (TWS), Apache Airflow, Azure DevOps
Developer Tools:	ServiceNow, Git, WinSCP, PuTTY
GenAI:	LangChain, LangGraph, Streamlit, ChromaDB
Familiar:	Django (Rest API), Oracle, MySQL, PostgreSQL

EXPERIENCE

TATA CONSULTANCY SERVICES   KAISER PERMANENTE (Mar 2023 – Present)

Data Engineer / AI CoE SPOC   *Hyderabad, India*

- Delivered **20TB+ of healthcare data** into **Azure Data Lake Storage (ADLS)** by **designing and optimizing enterprise-scale data pipelines** using **Azure Data Factory (ADF)** and **Databricks Delta Lake**, enabling faster analytics and reporting.
- Increased data processing efficiency by **60%** by **automating ETL/ELT workflows** across **ADF, ADLS, Databricks, and PySpark**, ensuring near **real-time data availability**.
- Improved data accuracy by **100%** through **implementing robust validation and quality frameworks**, delivering highly reliable and actionable business insights.
- Enabled seamless enterprise-wide analytics by **integrating on-premise, third-party, and Hadoop data sources** (**Oracle, Taleo, JSON, Avro, CSV, Parquet**) into **ADLS** via **ADF connectors**.
- Reduced Spark job runtime by **20%** for **2M+ daily records** by **implementing Spark performance tuning techniques** (caching, partitioning, broadcast joins) across **Hadoop environments**.
- Optimized **on-premise Hadoop clusters (HDFS, Hive, Spark)** by **fine-tuning resource allocation and data partitioning strategies**, improving large-scale batch processing performance by **35%** and ensuring smooth integration with **Azure cloud migration pipelines**.
- Serving as **SPOC (Single Point of Contact)** for the **AI CoE (Center of Excellence)**, building and integrating **AI agents** to automate workflows, enhance data-driven decision-making, and drive **AI innovation** across project teams.

TATA CONSULTANCY SERVICES   WALGREENS (Oct 2021 – Feb 2023)

Data Engineer   *Noida, India*

- Accelerated large-scale text processing by **90%** (2 hours → 15 minutes) by **developing and optimizing MapReduce programs** in the **Hadoop ecosystem**, significantly improving ETL throughput.
- Enhanced operational efficiency by **50%** by **automating data workflows** using **Python**, **HiveQL**, and **Shell scripting**, reducing manual intervention and improving reliability.
- Improved analytical data trustworthiness by **30%** through **data quality, consistency, and reconciliation checks** using **Hive**, **Impala**, and **Spark**.
- Streamlined ETL orchestration and reduced manual efforts by **40%** by **automating complex data pipelines** with **Python**, **HiveQL**, and **Shell**, achieving a **60% improvement in throughput**.
- Supported cloud migration initiatives by **managing data ingestion and batch processing** in **Hadoop Distributed File System (HDFS)**, laying the foundation for **Azure-based architectures**.
- Delivered enterprise-wide reporting solutions by **collaborating with business teams** and **processing large-scale datasets** across distributed ETL systems.

CERTIFICATIONS

Generative AI (Udemy) | 2025   [View Certificate ↗](#)

Databricks Fundamentals | 2025   [View Certificate ↗](#)

Azure Fundamentals (TCS Digital) | 2025   [View Certificate ↗](#)

Python Programming (TCS Digital) | 2022   [View Certificate ↗](#)

Microsoft Certified: Azure Data Fundamentals (DP-900) | Expected 2025   ⧗

OPEN SOURCE CONTRIBUTIONS

Liu Embeddings   [PyPI ↗] [Docs ↗]

Developed and published an open-source Python library, **liuembeddings**, which combines **text embeddings** and **vector storage** for efficient local NLP retrieval tasks. Built on **Hugging Face Transformers** and **ChromaDB**, it offers a **free, fully local, and privacy-friendly** alternative to cloud-based embedding services ideal for **AI projects** requiring fast and cost-efficient **embedding, storage, and retrieval** workflows.

PROJECTS

LeetCode SQL to PySpark (Top 50 Questions)   [View on GitHub ↗](#)

DataLemur SQL to PySpark   [View on GitHub ↗](#)

DataLemur SQL to PySpark (RDD Implementation)   [View on GitHub ↗](#)

Databricks Ecommerce analytics   [View on GitHub ↗](#)

Databricks ETL   [View on GitHub ↗](#)

AI Chatbot   [View on GitHub ↗](#)

The Analysis Bot   [View on GitHub ↗](#)

EDUCATION

NIET   Greater Noida, UP
Bachelor of Technology in Computer Science & Engineering   *Aug 2017 – Sep 2021*

ABC Public School   Gorakhpur, UP
Senior Secondary School   *Apr 2015 – Mar 2017*