

Himanshu Singh

+91 9129491699 | dev88himanshu@gmail.com

[LinkedIn](#) | [Portfolio](#)

PROFESSIONAL SUMMARY

Azure Certified Data Engineer with 4+ years of experience designing and optimizing **data pipelines**, **lakehouse architectures**, and **data warehouse solutions** across **Azure** and **on-premise Hadoop ecosystems**. Skilled in **ADF, Databricks, Spark, Hive, SQL, and Python**, delivering up to **60% performance improvements** and ensuring **data reliability at scale** for **Fortune 100 Healthcare and Retail (Pharmacy) clients**.

TECHNICAL SKILLS

Cloud & Data Engineering:	ADF, Databricks, Azure Synapse, ADLS/Blob Storage, Data Lakehouse
Big Data Ecosystem:	Spark (PySpark, RDD, Streaming), Hadoop (Hive, Impala, Sqoop, MapReduce, YARN)
Pipelines & Warehousing:	ETL/ELT, Data Modelling, Data Warehousing OLAP/OLTP
Programming:	Python, SQL, Shell Scripting
Orchestration & Workflow:	Tivoli Workload Scheduler (TWS), Apache Airflow, Azure DevOps
Developer Tools:	ServiceNow, Git, WinSCP, PuTTY
GenAI:	LangChain, LangGraph, Streamlit, ChromaDB
Familiar:	Django (Rest API), Oracle, MySQL, PostgreSQL

EXPERIENCE

TATA CONSULTANCY SERVICES

KAISER PERMANENTE (Mar 2023 – Present)

Data Engineer / AI CoE SPOC

Hyderabad, India

- Delivered **20TB+** of healthcare data into **Azure Data Lake Storage (ADLS)** by **designing and optimizing enterprise-scale data pipelines** using **Azure Data Factory (ADF)** and **Databricks Delta Lake**, enabling faster analytics and reporting.
- Increased data processing efficiency by **60%** by **automating ETL/ELT workflows** across **ADF, ADLS, Databricks, and PySpark**, ensuring near **real-time data availability**.
- Improved data accuracy by **100%** through **implementing robust validation and quality frameworks**, delivering highly reliable and actionable business insights.
- Enabled seamless enterprise-wide analytics by **integrating on-premise, third-party, and Hadoop data sources (Oracle, Taleo, JSON, Avro, CSV, Parquet)** into ADLS via **ADF connectors**.
- Reduced Spark job runtime by **20%** for **2M+ daily records** by **implementing Spark performance tuning techniques** (caching, partitioning, broadcast joins) across **Hadoop environments**.
- Optimized **on-premise Hadoop clusters (HDFS, Hive, Spark)** by **fine-tuning resource allocation and data partitioning strategies**, improving large-scale batch processing performance by **35%** and ensuring smooth integration with **Azure cloud migration pipelines**.
- Serving as **SPOC (Single Point of Contact)** for the **AI CoE (Center of Excellence)**, building and integrating **AI agents** to automate workflows, enhance data-driven decision-making, and drive **AI innovation** across project teams.

TATA CONSULTANCY SERVICES

WALGREENS (Oct 2021 – Feb 2023)

Data Engineer

Noida, India

- Accelerated large-scale text processing by **90%** (2 hours → 15 minutes) by **developing and optimizing MapReduce programs** in the **Hadoop ecosystem**, significantly improving ETL throughput.
- Enhanced operational efficiency by **50%** by **automating data workflows** using **Python, HiveQL, and Shell scripting**, reducing manual intervention and improving reliability.
- Improved analytical data trustworthiness by **30%** through **data quality, consistency, and reconciliation checks** using **Hive, Impala, and Spark**.
- Streamlined ETL orchestration and reduced manual efforts by **40%** by **automating complex data pipelines** with **Python, HiveQL, and Shell**, achieving a **60% improvement in throughput**.
- Supported cloud migration initiatives by **managing data ingestion and batch processing** in **Hadoop Distributed File System (HDFS)**, laying the foundation for **Azure-based architectures**.
- Delivered enterprise-wide reporting solutions by **collaborating with business teams** and **processing large-scale datasets** across distributed ETL systems.

OPEN SOURCE CONTRIBUTIONS

Liu Embeddings

[\[PyPI\]](#) [\[Docs\]](#)

Developed and published an open-source Python library, **liuembeddings**, which combines **text embeddings** and **vector storage** for efficient local NLP retrieval tasks. Built on **Hugging Face Transformers** and **ChromaDB**, it offers a **free, fully local, and privacy-friendly** alternative to cloud-based embedding services ideal for **AI projects** requiring fast and cost-efficient **embedding, storage, and retrieval** workflows.

DbMeta

[\[PyPI\]](#) [\[Git\]](#)

Developed and published an open-source, zero-dependency Python engine, **DbMeta**, designed for parsing and analyzing batch job metadata from **Hadoop, Spark, Databricks** by using folder-based job-run structures. Automatically detects **metadata** files, extracts structured metadata, and converts them into **queryable tables** for efficient monitoring and analytics.

CERTIFICATIONS

Databricks Fundamentals 2025	View Certificate
Databricks Certified Data Engineer Associate 2025	View Certificate
Azure Fundamentals (TCS Digital) 2025	View Certificate
Generative AI (Udemy) 2025	View Certificate
Python Programming (TCS Digital) 2022	View Certificate

PROJECTS

LeetCode SQL to PySpark (Top 50 Questions)	View on GitHub
DataLemur SQL to PySpark	View on GitHub
DataLemur SQL to PySpark (RDD Implementation)	View on GitHub
Databricks Ecommerce analytics	View on GitHub
Databricks ETL	View on GitHub
AI Chatbot	View on GitHub
The Analysis Bot	View on GitHub

EDUCATION

NIET

Bachelor of Technology in Computer Science & Engineering

Greater Noida, UP
Aug 2017 – Sep 2021

ABC Public School

Senior Secondary School

Gorakhpur, UP
Apr 2015 – Mar 2017